Zomato Personalised Food Recommendation

AI-driven recommendation concept to improve discovery, retention, and repeat orders

Product concept: "What's your Mood?"

Product Management Case Study · Independent concept

1. Executive Summary

Zomato has built one of India's stickiest consumer habits — roughly 21 million people order every month at an average of about 4.5 orders per user per month. But in early 2025 that habit stopped compounding: food-delivery growth flattened quarter-on-quarter, and the company itself attributed the softness to weak demand and rising competition. When new-user growth slows, the sharpest lever left is *existing* users ordering a little more often.

The obstacle sits at the very top of the journey. Zomato's discovery experience is built for *browsing* — an endless scroll of restaurants, filters and offers that assumes the user already knows what they want. Yet a large share of Indian users open the app hungry but undecided ("kuch samajh nahi aa raha kya khaayein"). For them, unlimited choice is not a feature — it is a tax. They scroll, tire, second-guess, and a meaningful fraction abandon *before they ever build a cart*.

"What's your Mood?" is an optional **express lane** that carries the undecided user from "I'm hungry" to "I've ordered" in seconds — without touching the browse feed that drives impulse buying and ad revenue. It is one product expressed through three surfaces:

- **Tell** — a natural-language search that understands intent ("something cold and sweet, not ice cream")
- **Offer** — mood tiles for users who can't put it into words
- **Remember** — *Fav Adda*, your usual orders surfaced proactively at the right moment

The concept enhances three surfaces Zomato already owns (search, home feed, reorder) rather than inventing new ones — keeping it feasible, incremental, and safe for the existing business. **North-star metric: repeat order frequency.**

2. Business Context

India's online food-delivery market is large and still growing fast — third-party estimates place it somewhere between roughly \$45B and \$55B in 2025, expanding at around 22% CAGR (estimates vary meaningfully by research firm, so this is directional). It is effectively a **duopoly**: Zomato and Swiggy together handle over 95% of restaurant-delivery orders, with Zomato leading at around 60% share.

Zomato's own reported numbers frame the opportunity precisely:

Metric	Figure	Why it matters
Monthly transacting customers (food delivery)	∼20.9M (Q4FY25)	The base every 0.1-order lift multiplies across
Order frequency	∼4.5 / user / month	Frequency *is* the business — the north star already exists
Food-delivery GOV	Dipped slightly QoQ (Q4FY25)	Growth is not guaranteed; the pressure is real
Recovery	Food-delivery NOV +18.8% YoY (Q4FY26)	Momentum returned, but competition stays intense

Two contextual facts shape the strategy. First, Zomato's parent rebranded to **Eternal Limited** in 2025 — the food-delivery app remains "Zomato". Second, the competitive clock is ticking: in September 2025 Swiggy launched an AI feature that predicts customer preferences from past orders to lift frequency. Intent-driven discovery is therefore not speculative — it is where the category leader's closest rival is already moving. The differentiation has to come from a *sharper* angle than "predict from history", which is exactly what a mood- and intent-led layer provides.

Sources: Eternal/Zomato quarterly shareholder letters (Q4FY25, Q4FY26), eternal.com investor relations; Business Standard and Indian Startup News (May 2025); Expert Market Research and IMARC/Renub market reports (2025–26); Market Research Future (Sept 2025). Company figures are from Eternal's audited shareholder letters; third-party market-size figures vary by source and are quoted as a range.

3. Problem Statement

Zomato's discovery experience optimises for users who **already know what they want**. It underserves the large segment who open the app hungry but undecided.

The friction is concentrated at a specific point in the funnel: **pre-cart decision fatigue**. The undecided user is met with hundreds of options and asked to do all the deciding. They browse for minutes, feel overwhelmed, and a meaningful share drop off *before a cart is ever created*.

This is deliberately scoped as a **discovery problem, not a checkout problem**. The abandonment we attack happens *before* the cart — driven by choice overload, not price shock or a broken payment screen. (Payment-stage abandonment is a separate, price-and-trust problem.) Precision keeps the solution, the metric and the success criteria clean.

Crucially, the users most hurt by this friction are the ones who matter most to the business: **returning users**, who open the app frequently — often out of routine rather than a specific craving — and who represent the bulk of repeat-order potential.

3.1 Goals & North Star

North-Star Metric: Repeat order frequency — average orders per returning user per month.

Chosen over acquisition or installs because it maps most directly to revenue (frequency is already the engine, so small lifts compound across ∼21M users); it is the *honest* metric for a decision-friction problem; and it resists vanity inflation from discounts or one-time signups.

Supporting metrics (leading indicators the north star is earned, not gamed):

- **Discovery efficiency** — time-to-cart and taps-before-first-add (should fall)
- **Session-to-order conversion** — % of app opens ending in an order

- **Feature-attributed orders** — % of orders via "What's your Mood?" surfaces vs. cold browsing (isolates our contribution)

Guardrail metric (what we must not break): **average order value & catalogue diversity** — a narrowing engine could lift frequency while shrinking baskets or trapping users in a loop. We watch AOV and cuisine variety to confirm we are deepening the habit, not starving it.

4. Research

Secondary research (completed). Public financials and market reports establish the business case: a flattening-then-recovering food-delivery segment, a frequency-driven model (∼4.5 orders/month), a duopoly under competitive pressure, and a market leader's rival already investing in AI preference prediction. Together these confirm that squeezing more repeat orders from existing users is a live priority, and that intent-driven discovery is a validated direction.

Primary research (proposed for validation). Before building, the following would de-risk the core assumption that decision fatigue causes measurable pre-cart drop-off:

- **Behavioural funnel analysis** — quantify the drop between app-open and first-cart-add; isolate "long browse, no order" sessions.
- **User interviews (8–12 returning users)** — probe the "hungry but undecided" moment and where they give up.
- **In-app survey** — a one-tap "why did you leave?" at abandonment to size "couldn't decide" vs. price vs. other.
- **Competitive teardown** — Swiggy's AI prediction feature and existing collections, to map where the intent gap still sits.

4.1 Target Users

Primary persona — the Returning User. Has order history, opens the app frequently, often without a specific dish in mind. Drives the repeat-order north star; every design decision is optimised for them.

Secondary persona — the New User. Little or no history; needs the app to reduce choice and build trust fast.

The two are not served by different products. The *same* three surfaces serve both — they reveal more as the app learns (graceful degradation):

Surface	**New user (no history)**	**Returning user (rich history)**
Tell (search)	Fully works — reads typed intent	Fully works
Offer (mood tiles)	Works with sensible default order	Re-ranks to their most-used moods
Remember (Fav Adda)	Hidden until first orders exist	Front and centre — core repeat driver

Persona = who we optimise for. Feature reach = who can use it. Optimise for the regular; lock nobody out.

4.2 Prioritisation — RICE

Several ideas surfaced during exploration. They were scored on **RICE** (Reach × Impact × Confidence ÷ Effort) to decide what makes v1 versus what is parked. Scores are relative and directional.

Idea	Reach	Impact	Conf.	Effort	Verdict
Fav Adda (proactive reorder)	High	High	High	Low	**Build (v1)**
Natural-language search	High	Med	Med	Med	**Build (v1)**
Mood tiles	High	Med	High	Low	**Build (v1)**
Pricing transparency	High	Med	High	Med	Park (separate trust initiative)
Home-chef ingredient kits	Med	High	Low	High	Park (Future — new business)
Gym / diet-only menus	Low	Med	Med	Med	Fold into "Guilt-free" tile

The three surfaces that scored highest — high reach, meaningful impact on the north star, and low-to-medium effort because they *enhance* existing surfaces — became "What's your Mood?". High-effort or lower-confidence ideas were deliberately parked, not discarded.

5. Insights

1. **Choice overload is the enemy of the undecided user.** More options help the decided shopper and paralyse the undecided one. The fix is not fewer options for everyone — it is a *faster path* for the stuck.
2. **The decisive moment is emotional and time-bound.** People don't want a "cuisine filter"; they want to answer "what am I in the mood for, right now?" Time, weather and day-of-week shape that answer.
3. **Reorder is the highest-intent, lowest-effort path to a repeat order — and it is under-exploited.** Existing reorder is a passive list users must hunt for. Surfaced *proactively at the right moment*, it becomes the strongest lever on the north star.
4. **Indian users prefer simple, intuitive interactions.** A blank "type what you want" box can freeze users. Adoption depends on low-effort entry: pre-filled prompts, one-tap tiles, and familiar, native language.

6. Hypothesis

Indian users prefer simple, intuitive interactions.

> **If** we give undecided users an optional express lane that translates their mood or intent into a short, relevant set of choices — and surfaces their usual orders proactively — **then** we will reduce pre-cart abandonment and increase repeat order frequency, **because** the friction causing drop-off is decision fatigue, not price or catalogue gaps.

This is falsifiable, which is the point. If discovery-efficiency metrics improve but repeat frequency does not, the hypothesis is wrong and the concept should be reconsidered — not quietly reframed.

7. Solution & Scope

7.1 The product: "What's your Mood?"

One intent engine, three surfaces, meeting the user wherever they are on the effort spectrum.

Tell — Natural-language search. The search bar understands plain language instead of keyword-matching. "Something cold and sweet but not ice cream" returns cold desserts, shakes and falooda — not restaurants named "Cool". Never shown empty: it rotates pre-filled, Hinglish-friendly example prompts so users learn what they can say. Core version is **one-shot** (type → understood → results); optional clarifying follow-ups are deferred to Future Enhancements. On a low-confidence parse it degrades to the closest mood collection plus a "did you mean…" nudge — never a dead "no results" screen.

Offer — Mood tiles. A compact, single-plane row: *Full-on hungry · Light bite · Something spicy · Something chill · Ghar ka khana · Guilt-free & high-protein.* One tap replaces 400 restaurants with a short, curated set. Tiles re-rank per user and per context (summer afternoon → "Something chill" leads; rainy evening → comfort leads).

Remember — Fav Adda. Your go-to orders from your go-to places, surfaced **proactively at the right moment** (open the app on a Friday evening → your usual biryani is already waiting, one tap to reorder). This is the primary lever on the north star. A one-line subtitle ("your go-to orders, one tap away") carries the meaning of the native name for first-time users.

7.2 Scope — what this is, and isn't

This concept **does not introduce standalone new features**. It enhances three surfaces Zomato already has:

- Search gains natural-language understanding.
- The home feed gains an optional mood-based fast path.
- Reorder evolves into a proactive *Fav Adda*.

The browse experience, ads, offers and full catalogue remain untouched. "What's your Mood?" is an *additive* express lane, not a redesign — it does not narrow the buffet, it adds a quick counter beside it. Every order it creates is therefore **incremental GMV**, recovered from a session heading toward zero, not GMV cannibalised from the impulse-driving scroll.

Explicitly out of scope (this version): ingredient/meal-kit delivery (a quick-commerce-shaped business — see Future Enhancements); pricing/checkout transparency (a separate trust initiative); conversational multi-turn search.

7.3 Why it drives repeat orders — the Hook Model

The concept is designed as a habit loop, mapping cleanly to Nir Eyal's **Hook Model**, which is why it targets the repeat-order north star rather than one-off conversion:

- **Trigger** — Fav Adda surfaces the user's usual proactively at the right moment (Friday evening → their biryani).
- **Action** — one tap to reorder, or one tap on a mood tile: the lowest-effort path to a cart.
- **Reward** — the user gets their comfort food fast, with the friction of deciding removed.
- **Investment** — every order teaches the engine (moods re-rank, Fav Adda grows), so the next loop is better — deepening the habit over time.

7.4 Competitive differentiation

The individual components are not novel — the value is in the combination and the framing. Zomato already has search, collections and a reorder button; Swiggy (Sept 2025) added AI that predicts preferences from past orders. The differentiation here is threefold: search understands *stated intent* in natural language rather than keyword-matching or history alone; reorder becomes *proactive anticipation* at the right moment rather than a passive list; and all three surfaces are unified into one intent engine framed as an *additive express lane* that protects — rather than cannibalises — the impulse-driven browse feed. It is a sharper angle than "predict from history", aimed squarely at repeat frequency.

8. Wireframes

Five screens illustrate the concept layered onto the existing app. Existing elements are kept as placeholders; only the three new surfaces are drawn.

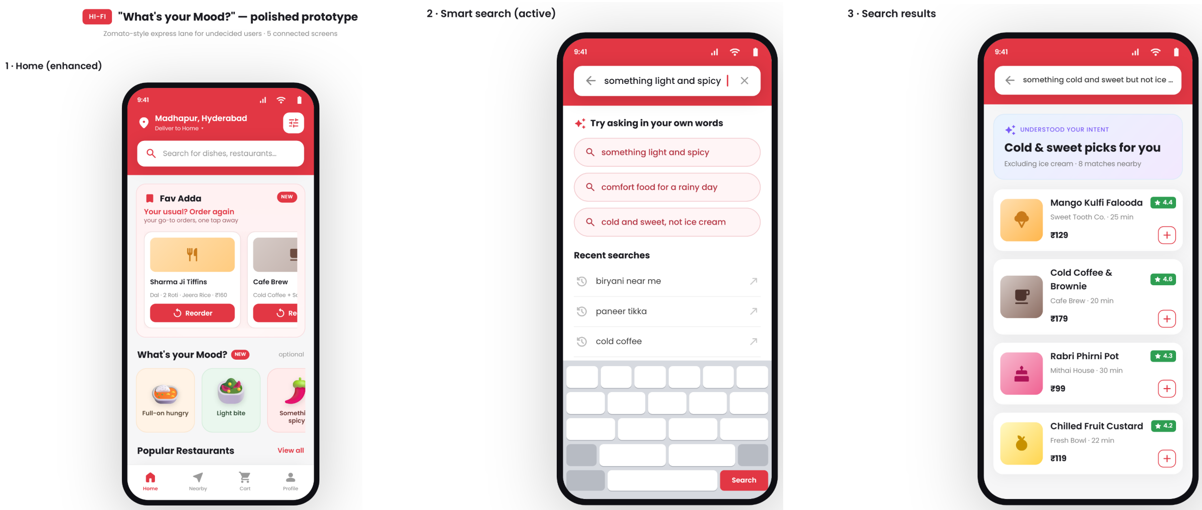

Figure 1: Left to right — (1) Home with Fav Adda + "What's your Mood?" row; (2) active smart search with intent prompts; (3) parsed search results echoing the understood intent.

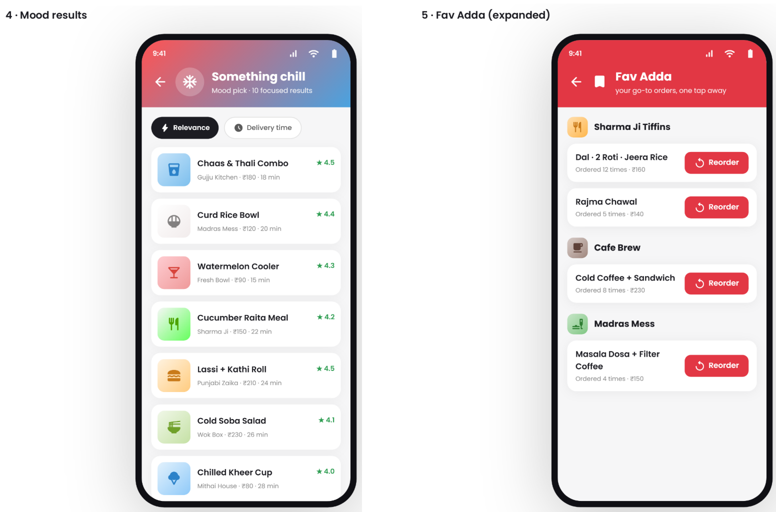

Figure 2: Left to right — (4) "Something chill" mood results, a deliberately short focused list; (5) expanded Fav Adda with usuals grouped by restaurant and one-tap reorder.

Screen notes. *S1 Home:* Fav Adda strip ("Your usual? Order again") + horizontal mood-tile row above the untouched restaurant list. *S2 Search:* focused bar with rotating Hinglish prompt chips (never empty). *S3 Results:* header echoes the understood intent. *S4 Mood:* ∼10 focused cards, not an endless list. *S5 Fav Adda:* usuals grouped by restaurant, one-tap reorder, clarifying subtitle.

9. User Journey

Current state (the pain):

Opens app (hungry, undecided) → endless scroll → applies a filter → scrolls → opens 3–4 restaurants → compares → *decision fatigue* → closes app without ordering.

The break happens between "scroll" and "cart". Nothing in the current flow helps the user *decide* — it only helps them *browse*.

Future state (with "What's your Mood?"):

> Opens app (hungry, undecided) → sees mood row + Fav Adda at top of home → **either** taps a mood tile (short curated set) **or** types intent into smart search **or** one-taps a Fav Adda usual → adds to cart → orders.

The undecided user now has three fast exits to an order; the decided user's browse experience is unchanged.

10. Success Metrics

Targets are illustrative and framed as hypotheses to validate via a controlled A/B experiment, not guarantees.

Tier	Metric	Direction / illustrative target
North star	Repeat order frequency (returning)	↑ from ∼4.5 toward ∼4.8–5.0 / user / month among adopters
Supporting	Session-to-order conversion	↑ (fewer zero-order sessions)
Supporting	Time-to-cart / taps-to-first-add	↓ (faster decisions)
Supporting	Feature-attributed orders	Establish baseline; grow share via the three surfaces
Supporting	Feature adoption	% of returning users engaging at least once/month
Guardrail	Average order value	Hold or ↑ (must not shrink)
Guardrail	Catalogue / cuisine diversity	Hold (must not collapse into a loop)

Approach: ship to a holdout-controlled cohort; read the north star over a 30–60 day window (frequency needs time to move); use feature-attributed orders to prove causation, not just correlation.

11. Trade-offs

- **Express lane vs. impulse browsing.** Guiding the undecided user could pull some sessions out of the impulse-heavy scroll. Mitigated by keeping the feature additive and the feed untouched; AOV is watched as a guardrail.
- **Personalisation vs. discovery diversity.** Surfacing "your usual" risks narrowing what users see. Guarded by the diversity metric and by keeping full browse one tap away.
- **Native naming vs. instant clarity ("Fav Adda").** Chose emotional, Indian-native branding over the clearer "My Regulars", accepting a small first-time-comprehension cost, mitigated by a subtitle.
- **One-shot vs. conversational search.** Chose one-shot for speed and to avoid "annoying bot" failure modes; conversational follow-ups deferred, not discarded.
- **Build effort vs. reuse.** Enhancing existing surfaces is more constrained than greenfield, but far more feasible and believable — a deliberate trade of ambition for shippability.

12. Future AI Enhancements

1. **Context-aware ranking.** Layer time-of-day, weather and day-of-week into ranking (rainy Chennai evening → pakora + chai rises). The clearest "feels like AI" upgrade.

2. **Reorder-cadence prediction.** Learn individual rhythms ("biryani most Fridays") and pre-empt them in Fav Adda — reorder from reactive to anticipatory.

3. **Conversational refinement.** Optional multi-turn follow-ups on ambiguous intent ("veg or non-veg?"), opt-in, never forced.

4. **Home-chef / meal-kit vision.** For the segment that misses home food but can't cook it, an ingredient-kit path (ingredients + steps + quantity-by-headcount). A separate, ambitious business — a north-facing vision, not a v1 commitment.

13. Conclusion

Zomato's growth is no longer guaranteed by new users, which makes existing-user frequency the sharpest available lever. The blocker is decision fatigue at the top of the journey — undecided users abandoning before they build a cart. **"What's your Mood?"** attacks exactly that: an optional express lane — smart search, mood tiles, and a proactive *Fav Adda* — that carries the stuck user from "I'm hungry" to "I've ordered", while leaving the browse feed, ads and impulse buying entirely intact.

It is feasible (it enhances surfaces that already exist), business-safe (every order it creates is incremental), and focused (one north star, one product, three surfaces). The differentiation is not in any single component — search, tiles and reorder all exist in some form — but in the *combination*: one intent engine, pointed squarely at repeat orders, framed to protect the very business it strengthens.

Company figures from Eternal Limited audited shareholder letters (FY25–FY26); market figures from third-party research firms, quoted as directional ranges. Independent concept, not affiliated with or endorsed by Zomato / Eternal Limited.